

The background features abstract, overlapping green geometric shapes, primarily triangles and polygons, in various shades of green, creating a modern and dynamic visual effect.

CS 35L

Software Construction Laboratory

Lecture 4.2

24th October, 2019

Logistics

- ▶ Assignment 4
 - ▶ Due on October 28th
- ▶ Hardware requirement for Week 8
 - ▶ Seeed Studio BeagleBone Green Wireless Development Board
- ▶ Assignment 10 Signup Sheet
 - ▶ Teams of 2
 - ▶ <https://docs.google.com/spreadsheets/d/1PVqVMEEsHjmmj9YLyqz5K4wU-k0Dwwm2iO9uS1Zq1wk/edit?usp=sharing>
 - ▶ Names must be filled by Friday of Week 3
 - ▶ Topic can be selected later
 - ▶ Ensure that topics are not the same
- ▶ Office Hours this week:
 - ▶ Friday - 9:30 - 11:30 am - BH3256S

Assignment 10 Rubric

- ▶ Presentation (50%):
 - ▶ Organization
 - ▶ Relevance to topic
 - ▶ Technical Details and Subject Knowledge
 - ▶ Presentation abilities (Elocution and Eye contact)
 - ▶ Content of slides
 - ▶ Ability to answer questions and interactivity with audience
- ▶ Report (50%)

Review - Previous Lab

- ▶ C Programming
 - ▶ Pointers, Double Pointers
 - ▶ Pass by Value, Pass by Reference
 - ▶ Structs
 - ▶ Dynamic Memory Allocation

Signed vs. Unsigned

<https://stackoverflow.com/questions/3812022/what-is-a-difference-between-unsigned-int-and-signed-int-in-c>

	two's complement	ones' complement	sign/magnitude
5	0000 0000 0000 0101	0000 0000 0000 0101	0000 0000 0000 0101
-5	1111 1111 1111 1011	1111 1111 1111 1010	1000 0000 0000 0101

- In two's complement, you get a negative of a number by inverting all bits then adding 1.
- In ones' complement, you get a negative of a number by inverting all bits.
- In sign/magnitude, the top bit is the sign so you just invert that to get the negative.

C Programming

Pointers to Functions

- ▶ A pointer that points to a function
- ▶ Declaration
 - ▶ `double (*func_ptr) (double, double);`
 - ▶ `func_ptr = &pow; // func_ptr points to pow()`
- ▶ Usage
 - ▶ `// Call the function referenced by func_ptr`
 - ▶ `double result = (*func_ptr)(1.5, 2.0);`
 - ▶ `// The same function call`
 - ▶ `result = func_ptr(1.5, 2.0);`

Function Pointers

- ▶ Variable which stores address to a function's executable code in memory.

```
#include <stdio.h>
void fun(int a)
{
    printf("Value of a is %d\n", a);
}
int main()
{
    void (*fun_ptr)(int) = &fun;
    (*fun_ptr)(10);
    return 0;
}
```


Qsort Example

```
#include <stdio.h>
#include <stdlib.h>

int compare (const void * a, const void * b) //qsort wants compare function to return int
{
    return ( *(int*)a - *(int*)b ); // typecasts void ptr to int ptr and then dereferences it
}

int main () {
    int values[] = { 40, 10, 100, 90, 20, 25 };
    qsort (values, 6, sizeof(int), compare); //pass compare function
    int n;
    for (n = 0; n < 6; n++) printf ("%d ",values[n]);
    return 0;
}
```

Formatted I/O

- ▶ `int fprintf(FILE * fp, const char * format, ...);`
- ▶ `int fscanf(FILE * fp, const char * format, ...);`
- ▶ `FILE *fp` can be either:
 - ▶ A file pointer
 - ▶ `stdin`, `stdout`, or `stderr`
- ▶ The format string
 - ▶ `int score = 120; char player[] = "Mary";`
 - ▶ `fp = fopen("file.txt", "w+")`
 - ▶ `fprintf(fp, "%s has %d points.\n", player, score);`

Debugging Process

- ▶ Reproduce the bug
- ▶ Simplify program input
- ▶ Use a debugger to track down the origin of the problem
- ▶ Fix the problem

Debugger

- ▶ A program that is used to run and debug other (target) programs
- ▶ Advantages:
 - ▶ Programmer can:
 - ▶ step through source code line by line
 - ▶ each line is executed on demand
 - ▶ interact with and inspect program at run-time
 - ▶ If program crashes, the debugger outputs where and why it crashed

GDB - GNU Debugger

- ▶ Debugger for several languages
 - ▶ C, C++, Java, Objective-C... more
- ▶ Allows you to inspect what the program is doing at a certain point during execution
- ▶ Semantic errors and segmentation faults are easier to find with the help of gdb

Using GDB

► Compile Program

- Normally: `$ gcc [flags] <source files> -o <output file>`
- Debugging: `$ gcc [other flags] -g <source files> -o <output file>`
 - enables built-in debugging support

► Specify Program to Debug

- `$ gdb <executable> or`
- `$ gdb`
- `(gdb) file <executable>`

Run-Time Errors

- ▶ Segmentation fault

- ▶ Program received signal SIGSEGV, Segmentation fault.
0x000000000400524 in function (arr=0x7fffc902a270, r1=2, c1=5,
r2=4, c2=6) at file.c:12
 - ▶ Line number where it crashed and parameters to the function that caused the error

- ▶ Semantic Error

- ▶ Program will run and exit successfully

- ▶ How do we find bugs?

Using GDB

▶ Run Program

- ▶ (gdb) run or
- ▶ (gdb) run [arguments]

▶ In GDB Interactive Shell

- ▶ Tab to Autocomplete, up-down arrows to recall history
- ▶ help [command] to get more info about a command

▶ Exit the gdb Debugger

- ▶ (gdb) quit

Factorial Program - to be Debugged

```
# include <stdio.h>
```

```
int main()
```

```
{
```

```
    int i, num, j;
```

```
    printf ("Enter the number: ");
```

```
    scanf ("%d", &num );
```

```
    for (i=1; i<num; i++)
```

```
        j=j*i;
```

```
    printf("The factorial of %d is %d\n",num,j);
```

```
}
```

Setting Breakpoints

► Breakpoints

- used to stop the running program at a specific point
- If the program reaches that location when running, it will pause and prompt you for another command

► Example:

- (gdb) break 6
 - Program will pause when it reaches line 6
- (gdb) break my_function
 - Program will pause at the first line of my_function every time it is called
- (gdb) break [position] if expression
 - Program will pause at specified position only when the expression evaluates to true

Breakpoints

- ▶ Setting a breakpoint and running the program will stop program where you tell it to
- ▶ You can set as many breakpoints as you want
 - ▶ (gdb) info breakpoints | break | br | b shows a list of all breakpoints

Basic Commands

- ▶ (gdb)step - Step to next line of code. Will step into a function.
- ▶ (gdb)next - Execute next line of code. Will not enter functions.
- ▶ (gdb)print <var> - Print value stored in variable.
- ▶ (gdb)continue - Continue execution to next break point.
- ▶ (gdb)set var <name>=<value> - Executes rest of program with new value of variable.

Deleting, Disabling and Ignoring BPs

- ▶ (gdb) delete [bp_number | range]
 - ▶ Deletes the specified breakpoint or range of breakpoints
- ▶ (gdb) disable [bp_number | range]
 - ▶ Temporarily deactivates a breakpoint or a range of breakpoints
- ▶ (gdb) enable [bp_number | range]
 - ▶ Restores disabled breakpoints
- ▶ If no arguments are provided to the above commands, all breakpoints are affected!!
- ▶ (gdb) ignore bp_number iterations
 - ▶ Instructs GDB to pass over a breakpoint without stopping a certain number of times.
 - ▶ bp_number: the number of a breakpoint
 - ▶ Iterations: the number of times you want it to be passed over

Displaying Data

- ▶ Why would we want to interrupt execution?
 - ▶ to see data of interest at run-time:
 - ▶ (gdb) print [/format] expression
 - ▶ Prints the value of the specified expression in the specified format
- ▶ Formats:
 - ▶ d: Decimal notation (default format for integers)
 - ▶ x: Hexadecimal notation
 - ▶ o: Octal notation
 - ▶ t: Binary notation

Resuming Execution After a Break

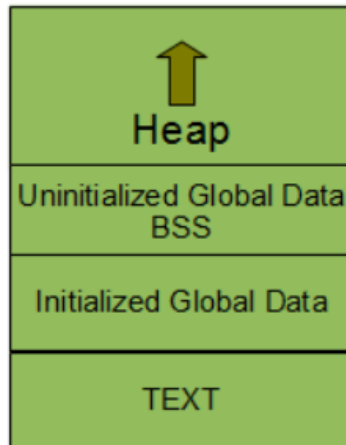

- ▶ When a program stops at a breakpoint
 - ▶ 4 possible kinds of gdb operations:
 - ▶ c or continue: debugger will continue executing until next breakpoint
 - ▶ s or step: debugger will continue to next source line
 - ▶ n or next: debugger will continue to next source line in the current (innermost) stack frame
 - ▶ f or finish: debugger will resume execution until the current function returns. Execution stops immediately after the program flow returns to the function's caller
 - ▶ the function's return value and the line containing the next statement are displayed

Watchpoints

- ▶ Watch/observe changes to variables
 - ▶ (gdb) watch my_var
 - ▶ sets a watchpoint on my_var
 - ▶ the debugger will stop the program when the value of my_var changes
 - ▶ old and new values will be printed
- ▶ (gdb) rwatch my_var
 - ▶ The debugger stops the program whenever the program reads the value of my_var

Process Memory Layout

(Higher Address)



(Lower Address)

- TEXT segment
 - Contains machine instructions to be executed
- Global Variables
 - Initialized
 - Uninitialized
- Heap segment
 - Dynamic memory allocation
 - malloc, free
- Stack segment
 - Push frame: Function invoked
 - Pop frame: Function returned
 - Stores
 - Local variables
 - Return address, registers, etc
- Command Line arguments and Environment Variables

Stack Info

- ▶ A program is made up of one or more functions which interact by calling each other
- ▶ Every time a function is called, an area of memory is set aside for it. This area of memory is called a stack frame and holds the following crucial info:
 - ▶ storage space for all the local variables
 - ▶ the memory address to return to when the called function returns
 - ▶ the arguments, or parameters, of the called function
- ▶ Each function call gets its own stack frame. Collectively, all the stack frames make up the call stack

Stack Frames and the Stack

```
1  #include <stdio.h>
2  void first_function(void);
3  void second_function(int);
4
5  int main(void)
6  {
7      printf("hello world\n");
8      first_function();
9      printf("goodbye goodbye\n");
10
11     return 0;
12 }
13
14
15 void first_function(void)
16 {
17     int imidate = 3;
18     char broiled = 'c';
19     void *where_prohibited = NULL;
20
21     second_function(imidate);
22     imidate = 10;
23 }
24
25
26 void second_function(int a)
27 {
28     int b = a;
29 }
```

Frame for `main()`

Frame for `first_function()`

Return to `main()`, line 9

Storage space for an int

Storage space for a char

Storage space for a void *

Frame for `second_function()`:

Return to `first_function()`, line 22

Storage space for an int

Storage for the int parameter named `a`

Analyzing the Stack in GDB

- ▶ (gdb) backtrace | bt
 - ▶ Shows the call trace (the call stack)
 - ▶ Without function calls:
 - ▶ #0 main () at program.c:10
 - ▶ one frame on the stack, numbered 0, and it belongs to main()
 - ▶ After call to function display()
 - ▶ #0 display (z=5, zptr=0xbffffb34) at program.c:15
 - ▶ #1 0x08048455 in main () at program.c:10
 - ▶ Two stack frames: frame 1 belonging to main() and frame 0 belonging to display().
 - ▶ Each frame listing gives
 - ▶ the arguments to that function
 - ▶ the line number that's currently being executed within that frame

Analyzing the Stack

- ▶ (gdb) info frame
 - ▶ Displays information about the current stack frame, including its return address and saved register values
- ▶ (gdb) info locals
 - ▶ Lists the local variables of the function corresponding to the stack frame, with their current values
- ▶ (gdb) info args
 - ▶ List the argument values of the corresponding function call

Other Useful Commands

- ▶ (gdb) info functions
 - ▶ Lists all functions in the program
- ▶ (gdb) list
 - ▶ Lists source code lines around the current line

Additional Info on GDB

- ▶ Gdb [cheat sheet](#)
- ▶ Gdb command [tutorial](#) and [slides](#)
- ▶ Running gdb [with emacs](#)

Assignment 4 - Laboratory

- ▶ Download old version of coreutils with buggy ls program
 - ▶ Untar, configure, make
- ▶ Bug: `ls -t` mishandles files whose time stamps are very far in the past. It seems to act as if they are in the future
 - ▶ `$ tmp=$(mktemp -d)`
 - ▶ `$ cd $tmp`
 - ▶ `$ touch -d '1918-11-11 11:00 GMT' wwi-armistice`
 - ▶ `$ touch now`
 - ▶ `$ sleep 1`
 - ▶ `$ touch now1`
 - ▶ `$ ls -lt wwi-armistice now now1`
- ▶ Output:
 - ▶ `-rw-r--r-- 1 eggert eggert 0 Nov 11 1918 wwi-armistice`
 - ▶ `-rw-r--r-- 1 eggert eggert 0 Feb 5 15:57 now1`
 - ▶ `-rw-r--r-- 1 eggert eggert 0 Feb 5 15:57 now`
- ▶ `$ cd`
- ▶ `$ rm -fr $tmp`

Fix the Bug!

- ▶ Reproduce the Bug
 - ▶ Follow steps on lab web page
- ▶ Simplify input
 - ▶ Run ls with -l and -t options only
- ▶ Debug
 - ▶ Use gdb to figure out what's wrong
 - ▶ `$ gdb ./ls`
 - ▶ `(gdb) run -lt /tmp/wwi-armistice /tmp/now /tmp/now1`
(run from the directory where the compiled ls lives)
- ▶ Patch
 - ▶ Construct a patch "lab4.diff" containing your fix
 - ▶ It should contain a ChangeLog entry followed by the output of `diff -u`

Hints

- ▶ Don't forget to answer all questions! (lab4.txt)
- ▶ Make sure not to submit a reverse patch! (lab4.diff)
- ▶ “Try to reproduce the problem in your home directory, instead of the \$tmp directory. How well does SEASnet do?”
 - ▶ Timestamps represented as seconds since Unix Epoch
 - ▶ SEASnet NFS filesystem has unsigned 32-bit time stamps
 - ▶ Local File System on Linux server has signed 32-bit time stamps
 - ▶ If you touch the files on the NFS filesystem it will return timestamp around 2054
 - ▶ => files have to be touched on local filesystem (df -l)
- ▶ Use “info functions” to look for relevant starting point
- ▶ Use “info locals” to check values of local variables

Assignment 10 - Presentations

- ▶ Today's Presentation

- ▶ Alvin Nguyen
- ▶ Henry Kou

- ▶ Next Class

- ▶ Sheet not filled!

Questions?